



# Distributed Systems

Lab Introduction  
Part 2

Charalampos Stylianopoulos

# Agenda

- Lab 2 introduction
- Solution costs for Lab 1
- Concurrency in Repy

# Notes about the labs

- Lab rooms:
  - Prefer rooms 3354 and 3358
  - That's where TAs will be.
- Check ping pong, lab pages for news, FAQs etc.
  - E.g. How to use curl

```
for i in `seq 1 20`; do
curl -d 'entry=t'${i} -X 'POST' 'http://ip:port/entries'
done
```

This will post entry=t1, entry=t2, ..., entry=t20 to <http://ip:port/entries>

# Notes about the labs(2)

- We expect you to use the RESTful API for browser –vessel communication.
- You also want to extend it and use it for vessel-vessel communication
  - It will make your life easier when the code becomes more complex

Functions	API	Parameters	Returns
Add a new entry	POST /entries	entry : text	Status
Modify an entry	PUT /entries/ entryID	entry : text	Status
Delete an entry	DELETE /entries/ entryID	None	Status
Modify or Delete an entry	POST /entries/ entryID	entry : text <i>delete: logical</i>	Status
Add entry to neighbor	POST /neighbourID/ entries	<i>entry:text</i> <i>neighbourID:text</i>	Status

Lab 2 Introduction

# **RELIABLE CENTRALIZED BLACKBOARD**

# Distributed Blackboard


- We have a simple working version so far...  
Let's make it better!
- Reliable and consistent
  - Every board shows messages in the same order
  - No message gets lost
- How? Centralized version!

# Distributed Blackboard – Centralized

- Each post is sent to the leader which distributes it to the network
- The leader should be able to handle correctly multiple posts from different nodes...
- But who is the leader?

# Leader election on a ring

1. Node 5 initiates the leader election process. It sends ID to its next node in an “Election” (E) message.
2. When node 6 receives the message, it appends its ID and forwards the message
3. When the message gets back to the process that started the election:
  - i. it elects process with highest ID as coordinator, and
  - ii. changes the message type to “Coordination” message (C) and circulates it in the ring





# Leader Election

- Use the Ring-based Election Algorithm (see Lecture slides on Monday) when starting the board in order to decide the leader
  - Define a ring topology
  - Every node should send **only** to their next neighbor
  - Use a **locally generated random number** as a criterion for selecting the leader (e.g. highest wins)
- Every node acts as an initiator in the beginning .
  - n elections running **concurrently**.
  - Eventually they all agree on the same leader.
- The protocol starts running as soon as the nodes are up .
  - you might have to wait a bit to make sure everyone has booted .
  - How? Use a timer (see slide later on).
- Simplifications (but feel free to impress us):
  - Not dynamic – only run election in the initialization of the protocol
  - Assume that communication between neighbors is reliable

# After the election

- After the election, nodes send messages directly to the leader.
- The leader can serve as centralized sequencer:
  - He decides the correct, global order of all messages.
  - Everybody else follows that order.
  - Hint: use the sequence number field. It will now be issued by the leader and it will be unique.

# What could go wrong?

Ensure that the ordering of the blackboard entries is the same on every vessel, **even in corner cases.**

Pitfalls:

- **waitforconn(ip,port,function)**, spawns a new thread to run **function**
- if **function** accesses shared variables or data structures, you need to take care of any synchronization needed (i.e. no thread safety when accessing shared data)

# What could go wrong?(2)

Possible race conditions:

- On the Non-Leaders: Multiple browsers operating on the same blackboard.
- On the leader: Messages from different blackboards will be handled by different threads.

Hint: Use a lock every time you read/write to a shared resource.

# Task 1

## Leader Election

- Explain your leader election algorithm
- Use a field in the webpage to show who the leader is and what its random number is
- Discuss the *solution cost* of the leader election algorithm that you use\*

\*= slide 21

# Task 2

## Blackboard (centralized)

- Show that concurrent **submissions** do not lead to problems anymore
  - with multiple browsers submitting in the same vessel concurrently, and
  - with multiple vessels submitting concurrently
- Explain where you use locks and why (in case you do use locks)
- Demonstrate the cost of your solution (i.e. cost of a post delivered to all nodes)
- Briefly discuss pros + cons of this design

# Optional Tasks

- Note: completely optional
  - We still give you up to 10 points even without this extension
- Handle dynamic networks:
  - What happens if the leader fails while the program is running?
  - What happens if a node during the election cannot reach its next neighbor?
- Concurrently delete/modify entries in the blackboard.

# Summary

- ✓ Leader election protocol to decide a leader among the blackboards.
- ✓ Blackboard must now be consistent always.

Optional:

- Dynamic leader election.
- Delete/Modify.
- ✓ Deliverables: code + Video(or report).

Deadline: November 24




# Agenda

- Lab 2 introduction
- **Solution costs for Lab 1**
- Concurrency in Repy



# Solution cost

- We can measure the cost of a solution (e.g. in Lab 1) in terms of
  - **number of nodes** to which a new post is propagated
  - **payload**: number of blackboard entries per message
- Overall cost per post =  $(\text{number\_of\_nodes} - 1) \cdot (\text{payload})$
- For example, consider the case of three vessels A, B and C, and the following events:
  - Event 1: User posts “msg1” to vessel A.
  - Event 2: User posts “msg2” to vessel A.





# Example: a good scenario

- Upon Event 1,  
Vessel A sends “msg1” to vessels B and C
  - Payload for each message = 1
  - Overall cost = 2
- Upon Event 2,  
Vessel A sends “msg2” to vessels B and C
  - Payload for each message = 1
  - Overall cost = 2
- Overall:  $m(n-1)$



# Example: a costly scenario

- Upon Event 1, Vessel A sends “msg1” to vessels B and C
  - Payload for each message = 1
  - Overall cost = 2
- Upon Event 2, Vessel A sends “msg1, msg2” to vessels B and C
  - Payload for each message = 2
  - Overall cost = 4
- Overall:  $bm(n-1)$ 
  - board size \* msg \* (#nodes-1)



# Cost of Lab2

- We want you do the same simple communication cost analysis for Lab2:
  - First for the leader election part.
  - Then for the blackboard.

# Agenda

- Lab 2 introduction
- Solution costs for Lab 1
- **Concurrency in Repy**

# Concurrency in Rpy

- Locks
- Example 1
- Timers
- Example 2

# Locks

- **getlock()**  
Returns a lock object that can be used for mutual exclusion and critical section protection.
- **lock.acquire(blocking=1)**  
Blocks until the lock is available, then takes it (lock is an object obtained by calling getlock()).
  - If the optional "blocking" argument is False, the method returns False immediately instead of waiting to acquire the lock; if the lock is available it takes it and returns True, as if it were called with no argument.
- **lock.release()**  
Releases the lock. Do not call it if the lock is unlocked.



# Lock example

Write to a log file on each request, and make sure that the file is only handled by **one thread at a time**.

```
def on_request(ip, port, socket, handle, listener):  
    mycontext['lock'].acquire()  
  
    # Here's the critical section:  
    log_file = open('log.txt', 'a')  
    log_file.write('got a request from ' + ip + '\n')  
    log_file.close()  
  
    mycontext['lock'].release()  
  
if callfunc == 'initialize':  
    mycontext['lock'] = getlock()  
    waitforconn(getmyip(), 63153, on_request)
```

# Lock example

Write to a log file on each request, and make sure that the file is only handled by **one thread at a time**.

```
def on_request(ip, port, socket, handle, listener):  
    mycontext['lock'].acquire()
```

```
    # Here's the critical section:  
    log_file = open('log.txt', 'a')  
    log_file.write('got a request from ' + ip + '\n')  
    log_file.close()
```

Critical  
Section

```
    mycontext['lock'].release()
```

```
if callfunc == 'initialize':  
    mycontext['lock'] = getlock()  
    waitforconn(getmyip(), 63153, on_request)
```

# Lock example

Write to a log file on each request, and make sure that the file is only handled by **one thread at a time**.

```
def on_request(ip, port, socket, handle, listener):  
    mycontext['lock'].acquire()  
  
    # Here's the critical section:  
    log_file = open('log.txt', 'a')  
    log_file.write('got a request from ' + ip + '\n')  
    log_file.close()  
  
    mycontext['lock'].release()  
  
if callfunc == 'initialize':  
    mycontext['lock'] = getlock()  
    waitforconn(getmyip(), 63153, on_request)
```



Lock initialization.

# Lock example

Write to a log file on each request, and make sure that the file is only handled by **one thread at a time**.

```
def on_request(ip, port, socket, handle, listener):  
    mycontext['lock'].acquire()  
  
    # Here's the critical section:  
    log_file = open('log.txt', 'a')  
    log_file.write('got a request from ' + ip + '\n')  
    log_file.close()  
  
    mycontext['lock'].release()  
  
if callfunc == 'initialize':  
    mycontext['lock'] = getlock()  
    waitforconn(getmyip(), 63153, on_request)
```

Will spawn a new thread on a new connection...



# Lock example

Write to a log file on each request, and make sure that the file is only handled by **one thread at a time**.

```
def on_request(ip, port, socket, handle, listener):  
    mycontext['lock'].acquire()  
  
    # Here's the critical section:  
    log_file = open('log.txt', 'a')  
    log_file.write('got a request from ' + ip + '\n')  
    log_file.close()  
  
    mycontext['lock'].release()  
  
if callfunc == 'initialize':  
    mycontext['lock'] = getlock()  
    waitforconn(getmyip(), 63153, on_request)
```

..and every new  
thread will call this  
method

Will spawn a new thread on a  
new connection...

# Lock example

Write to a log file on each request, and make sure that the file is only handled by **one thread at a time**.

```
def on_request(ip, port, socket, handle, listener):
```

```
    mycontext['lock'].acquire()
```

```
    # Here's the critical section:
```

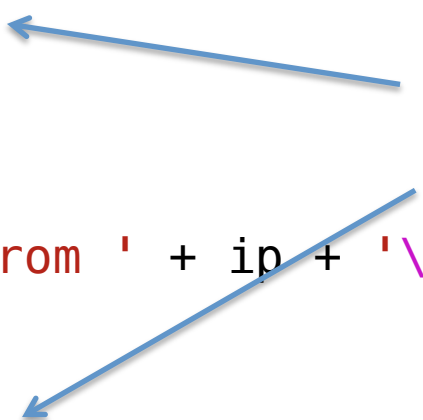
```
    log_file = open('log.txt', 'a')
```

```
    log_file.write('got a request from ' + ip + '\n')
```

```
    log_file.close()
```

```
    mycontext['lock'].release()
```

Acquire and  
release the  
lock to protect  
access



```
if callfunc == 'initialize':
```

```
    mycontext['lock'] = getlock()
```

```
    waitforconn(getmyip(), 63153, on_request)
```

# Timers

- **settimer(waittime, function, args)**  
Sets a timer that when it expires will start a new thread to call a function with a set of arguments.
- **canceltimer(timerhandle)**  
Tries to cancel a timer handle that has not started a thread.

# Timer example

Count the number of requests received, and print the result when no requests have been received for 10 seconds.

```
def on_request(ip, port, socket, handle, listener):
    mycontext['counter'] += 1
    canceltimer(mycontext['timer'])
    mycontext['timer'] = settimer(10, stop, [listener])

# This is the function that will run in a separate thread:
def stop(listener):
    print 'got ' + str(mycontext['counter']) + ' requests'
    stopcomm(listener)

if callfunc == 'initialize':
    mycontext['counter'] = 0
    listener = waitforconn(getmyip(), 63153, on_request)
    mycontext['timer'] = settimer(10, stop, [listener])
```



# Timer example

Count the number of requests received, and print the result when no requests have been received for 10 seconds.

```
def on_request(ip, port, socket, handle, listener):
    mycontext['counter'] += 1
    canceltimer(mycontext['timer'])
    mycontext['timer'] = settimer(10, stop, [listener])

# This is the function that will run in a separate thread:
def stop(listener):
    print 'got ' + str(mycontext['counter']) + ' requests'
    stopcomm(listener)

if callfunc == 'initialize':
    mycontext['counter'] = 0
    listener = waitforconn(getmyip(), 63153, on_request)
    mycontext['timer'] = settimer(10, stop, [listener])
```

Start a timer, that  
expires in 10s.

# Timer example

Count the number of requests received, and print the result when no requests have been received for 10 seconds.

```
def on_request(ip, port, socket, handle, listener):
    mycontext['counter'] += 1
    canceltimer(mycontext['timer'])
    mycontext['timer'] = settimer(10, stop, [listener])

# This is the function that will run in a separate thread:
def stop(listener):
    print 'got ' + str(mycontext['counter']) + ' requests'
    stopcomm(listener)

if callfunc == 'initialize':
    mycontext['counter'] = 0
    listener = waitforconn(getmyip(), 63153, on_request)
    mycontext['timer'] = settimer(10, stop, [listener])
```

If it expires, it spawns a thread that calls “stop”

Start a timer, that expires in 10s.

# Timer example

Count the number of requests received, and print the result when no requests have been received for 10 seconds.

```
def on_request(ip, port, socket, handle, listener):  
    mycontext['counter'] += 1  
    canceltimer(mycontext['timer'])  
    mycontext['timer'] = settimer(10, stop, [listener])
```

A new connection  
will try to reset the  
timer

# This is the function that will run in a separate thread:

```
def stop(listener):  
    print 'got ' + str(mycontext['counter']) + ' requests'  
    stopcomm(listener)
```

```
if callfunc == 'initialize':  
    mycontext['counter'] = 0  
    listener = waitforconn(getmyip(), 63153, on_request)  
    mycontext['timer'] = settimer(10, stop, [listener])
```

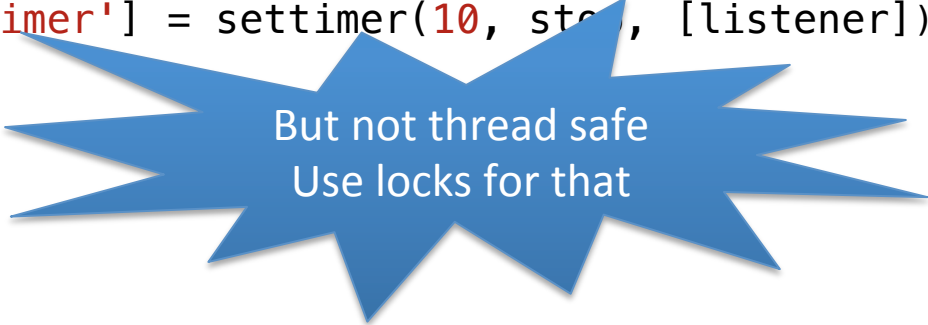
# Timer example

Count the number of requests received, and print the result when no requests have been received for 10 seconds.

```
def on_request(ip, port, socket, handle, listener):
    mycontext['counter'] += 1
    canceltimer(mycontext['timer'])
    mycontext['timer'] = settimer(10, stop, [listener])

# This is the function that will run in a separate thread:
def stop(listener):
    print 'got ' + str(mycontext['counter']) + ' requests'
    stopcomm(listener)

if callfunc == 'initialize':
    mycontext['counter'] = 0
    listener = waitforconn(getmyip(), 63153, on_request)
    mycontext['timer'] = settimer(10, stop, [listener])
```



But not thread safe  
Use locks for that

# More Examples

- <https://seattle.poly.edu/wiki/RepyTutorial#RacesSleepandLocksexample1.6>